

# iris

September 13, 2024

```
[211]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.metrics import (confusion_matrix , accuracy_score,
    ConfusionMatrixDisplay ,classification_report ,
precision_score,recall_score)
from sklearn.ensemble import (AdaBoostClassifier ,BaggingClassifier,
    GradientBoostingClassifier,
RandomForestClassifier)
from sklearn.svm import SVC
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
```

```
[11]: # load iris datasets
iris = pd.read_csv('IRIS.csv')
iris.head()
```

```
[11]:   sepal_length  sepal_width  petal_length  petal_width  species
0          5.1           3.5           1.4           0.2  Iris-setosa
1          4.9           3.0           1.4           0.2  Iris-setosa
2          4.7           3.2           1.3           0.2  Iris-setosa
3          4.6           3.1           1.5           0.2  Iris-setosa
4          5.0           3.6           1.4           0.2  Iris-setosa
```

```
[17]: iris.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
 #   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
```

```
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
[21]: iris.isna().sum()
```

```
[21]: sepal_length    0
      sepal_width    0
      petal_length   0
      petal_width    0
      species        0
      dtype: int64
```

```
[25]: iris.describe(include='all')
```

```
[25]:
```

	sepal_length	sepal_width	petal_length	petal_width	species
count	150.000000	150.000000	150.000000	150.000000	150
unique	NaN	NaN	NaN	NaN	3
top	NaN	NaN	NaN	NaN	Iris-setosa
freq	NaN	NaN	NaN	NaN	50
mean	5.843333	3.054000	3.758667	1.198667	NaN
std	0.828066	0.433594	1.764420	0.763161	NaN
min	4.300000	2.000000	1.000000	0.100000	NaN
25%	5.100000	2.800000	1.600000	0.300000	NaN
50%	5.800000	3.000000	4.350000	1.300000	NaN
75%	6.400000	3.300000	5.100000	1.800000	NaN
max	7.900000	4.400000	6.900000	2.500000	NaN

```
[27]: iris.nunique()
```

```
[27]: sepal_length    35
      sepal_width    23
      petal_length   43
      petal_width    22
      species        3
      dtype: int64
```

```
[54]: iris.iloc[:, :4].var()
```

```
[54]: sepal_length    0.685694
      sepal_width    0.188004
      petal_length    3.113179
      petal_width    0.582414
      dtype: float64
```

# 1 Modeling

```
[48]: iris.columns
```

```
[48]: Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',  
         'species'],  
        dtype='object')
```

```
[57]: # convert species column to dummies by labelencoder  
from sklearn.preprocessing import LabelEncoder  
label = LabelEncoder()  
label.fit_transform(iris['species'])
```

```
[57]: array([[0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
          0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,  
          0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
          1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
          1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
          2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
          2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
          2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2])
```

```
[59]: # split our data into features and targets  
X = iris.drop('species',axis=1)  
y = iris['species']
```

```
[61]: # train test split our data  
X_train,X_test,y_train,y_test = train_test_split(X,y,stratify=y,test_size=0.  
↪3,random_state=42)
```

```
[225]: # SVC -> Model # 1  
clf = SVC(C=10,kernel='rbf')  
clf.fit(X_train, y_train)  
y_pred = clf.predict(X_test)  
  
# Evaluation  
print(accuracy_score(y_test, y_pred))  
print(confusion_matrix(y_test, y_pred))  
print(ConfusionMatrixDisplay(confusion_matrix(y_test, y_↵  
↪y_pred),display_labels=clf.classes_).plot())  
print(classification_report(y_test,y_pred))  
print(recall_score(y_test,y_pred,average='micro'))  
print(precision_score(y_test,y_pred,average='macro'))
```


```
0.9777777777777777  
[[15  0  0]  
 [ 0 14  1]  
 [ 0  0 15]]
```

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x7f84f14632e0>

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	15
Iris-versicolor	1.00	0.93	0.97	15
Iris-virginica	0.94	1.00	0.97	15
accuracy			0.98	45
macro avg	0.98	0.98	0.98	45
weighted avg	0.98	0.98	0.98	45

0.9777777777777777

0.9791666666666666



```
[227]: # KNN -> Model # 2
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)
y_pred = knn.predict(X_test)

# Evaluation
print(accuracy_score(y_test, y_pred).round(2))
```

```

print()
from sklearn.metrics import confusion_matrix
print(confusion_matrix(y_test, y_pred))
print(ConfusionMatrixDisplay(confusion_matrix(y_test,
↪y_pred),display_labels=knn.classes_).plot())
print(classification_report(y_test,y_pred))
print(recall_score(y_test,y_pred,average='micro'))
print(precision_score(y_test,y_pred,average='macro'))

```

0.98

```

[[15  0  0]
 [ 0 15  0]
 [ 0  1 14]]


```

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x7f84f18f8670>

	precision	recall	f1-score	support
Iris-setosa	1.00	1.00	1.00	15
Iris-versicolor	0.94	1.00	0.97	15
Iris-virginica	1.00	0.93	0.97	15
accuracy			0.98	45
macro avg	0.98	0.98	0.98	45
weighted avg	0.98	0.98	0.98	45

0.9777777777777777

0.9791666666666666



```
[229]: # Logistic Regression -> Model # 3
log = LogisticRegression(max_iter=10000)
log.fit(X_train, y_train)
y_pred = log.predict(X_test)


# Evaluation
from sklearn.metrics import accuracy_score
print(accuracy_score(y_test, y_pred))
print()
from sklearn.metrics import confusion_matrix
print(confusion_matrix(y_test, y_pred))
print(ConfusionMatrixDisplay(confusion_matrix(y_test, y_pred),
    display_labels=log.classes_).plot())
print(recall_score(y_test, y_pred, average='micro'))
print(precision_score(y_test, y_pred, average='macro'))
```

0.9333333333333333

```
[[15  0  0]
 [ 0 14  1]
 [ 0  2 13]]
```

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at

```
0x7f84f18f8e50>  
0.9333333333333333  
0.9345238095238096
```



```
[231]: # Random Forests -> Model 4  
rfc = RandomForestClassifier(n_estimators=500,max_depth=50)  
rfc.fit(X_train, y_train)  
y_pred = rfc.predict(X_test)  
  
# Evaluation  
from sklearn.metrics import accuracy_score  
print(accuracy_score(y_test, y_pred))  
print()  
from sklearn.metrics import confusion_matrix  
print(confusion_matrix(y_test, y_pred))  
print(ConfusionMatrixDisplay(confusion_matrix(y_test, y_pred),  
    y_pred),display_labels=rfc.classes_.plot())  
print(recall_score(y_test,y_pred,average='micro'))  
print(precision_score(y_test,y_pred,average='macro'))
```


```
0.9111111111111111
```

```
[[15  0  0]
 [ 0 14  1]
 [ 0  3 12]]
```

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x7f84f2294940>

```
0.9111111111111111
```

```
0.9155354449472096
```



```
[233]: # GradientBoostingClassifier -> Model 5
grc = GradientBoostingClassifier()
grc.fit(X_train, y_train)
y_pred = grc.predict(X_test)

# Evaluation
print(accuracy_score(y_test, y_pred))
print()
print(confusion_matrix(y_test, y_pred))
print(ConfusionMatrixDisplay(confusion_matrix(y_test, y_
    ↪ y_pred), display_labels=grc.classes_).plot())
print(recall_score(y_test, y_pred, average='micro'))
print(precision_score(y_test, y_pred, average='macro'))
```



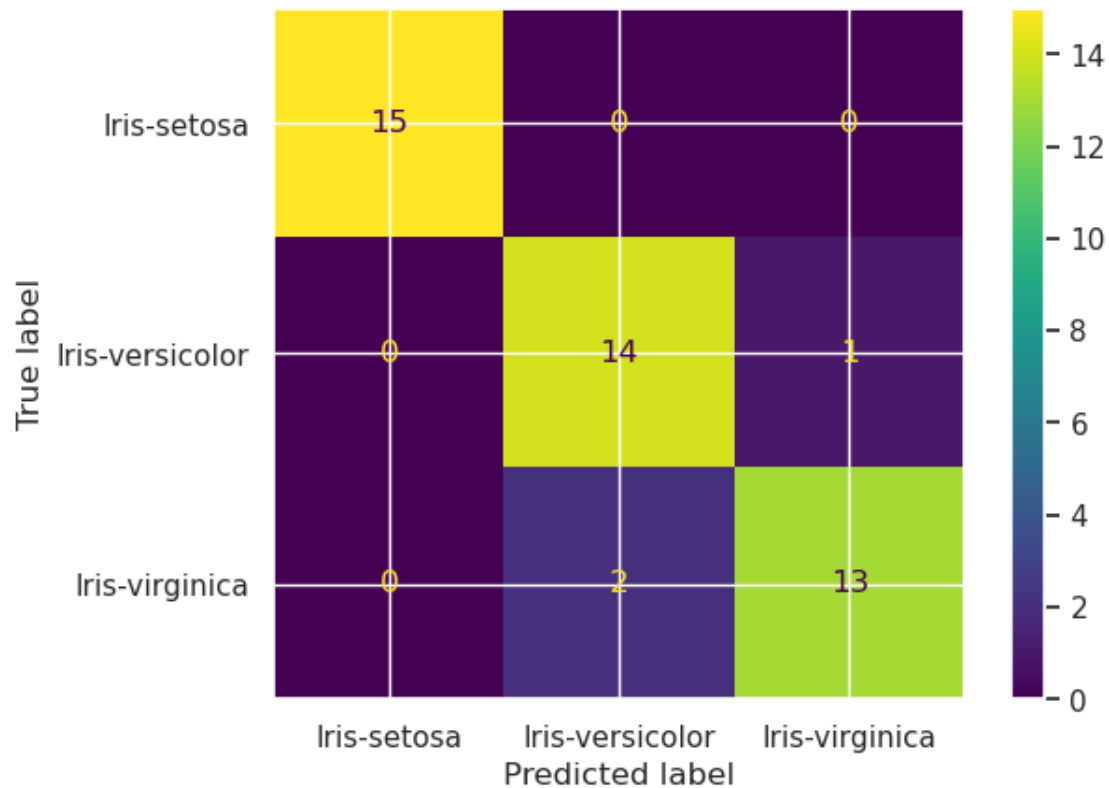
0.9333333333333333

```
[[15  0  0]
 [ 0 14  1]
 [ 0  2 13]]
```

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x7f84f22955d0>

0.9333333333333333

0.9345238095238096



```
[235]: # Gradient Boosting Classifier -> Model 6
ada = AdaBoostClassifier()
ada.fit(X_train, y_train)
y_pred = ada.predict(X_test)

# Evaluation
print(accuracy_score(y_test, y_pred))
print()
print(confusion_matrix(y_test, y_pred))
print(ConfusionMatrixDisplay(confusion_matrix(y_test, y_pred),
display_labels=ada.classes_).plot())
print(recall_score(y_test, y_pred, average='micro'))
```

```
print(precision_score(y_test,y_pred,average='macro'))
```


0.9333333333333333

```
[[15  0  0]
 [ 0 14  1]
 [ 0  2 13]]
```

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x7f84f260da80>

0.9333333333333333

0.9345238095238096



```
[239]: # Gradient Boosting Classifier -> Model 7
bag = BaggingClassifier(n_estimators=100)
bag.fit(X_train, y_train)
y_pred = bag.predict(X_test)

# Evaluation
print(accuracy_score(y_test, y_pred))
print()
print(confusion_matrix(y_test, y_pred))
```

```
print(ConfusionMatrixDisplay(confusion_matrix(y_test, y_
↪y_pred), display_labels=bag.classes_).plot())
print(recall_score(y_test, y_pred, average='micro'))
print(precision_score(y_test, y_pred, average='macro'))
```

0.9333333333333333

```
[[15  0  0]
 [ 0 14  1]
 [ 0  2 13]]
```

<sklearn.metrics.\_plot.confusion\_matrix.ConfusionMatrixDisplay object at 0x7f84f2570fa0>

0.9333333333333333

0.9345238095238096

